

Maintainer Interview Guide

GenAI Usage and Burden

Purpose

This guide supports one on one, recorded interviews with open source package maintainers. The goal is to understand how GenAI and LLM tools are affecting maintainer workload and burden, not just how often people use these tools. Findings will feed into a pyOpenSci writeup and reporting effort on overall maintainer burden.

How to use this guide

Treat this as semi-structured, not a script. The core section is designed to run about 30 minutes. If the conversation is going well and there is time, move into the extension section, which can extend the conversation to about 60 minutes. Follow interesting threads even if they mean skipping ahead or out of order. Probes are suggestions, not requirements. Use them when a question lands flat or when there is room to go deeper.

Intro and consent (about 2 minutes)


Read or paraphrase:

“Thanks for making time for this. This is part of a pyOpenSci effort looking at how GenAI tools are affecting package maintainers, specifically around workload and burden, not just adoption rates. I would like to record this so I can stay in the conversation instead of taking notes. The recording will be used for transcription and analysis, and any published reporting will anonymize your responses. Is that okay with you?”

We have about 30 minutes, maybe up to 60 if it is flowing and you have time. There are no wrong answers here. I am interested in your honest experience, including the parts that feel messy or contradictory.”

Stay involved

Want to stay involved after this interview? Scan the code or use the link below to join the survey, sprint, or focus group mailing list.



QR CODE
PLACEHOLDER

Survey link: _____

Core block (about 30 minutes)

Context and workflow (about 5 minutes)

- What project or projects do you maintain, and how long have you been doing it?
- How big is your maintainer team? How much of this lands on you alone?
- Where does GenAI show up in your day to day maintenance work right now, if at all? Think coding, review, triage, docs, communications.

Probe: if they say GenAI does not show up in their own work, ask whether they see it showing up in others' contributions even so.

Motivations and perceptions (about 7 minutes)

- What led you to start using GenAI tools in your maintainer work? Or, what has kept you from adopting them?
- When you do use them, what are you actually gaining? Not how often you use them, but what the value is.
- What are your honest feelings about these tools in your role as a maintainer specifically?
- Has that perception shifted over time? Better, worse, more mixed?

Probe: is there a gap between what you think you are supposed to feel about these tools and what you actually feel day to day?

Burden and strain (about 8 minutes)

- Where has GenAI added burden to your work? Where has it relieved it?
- Walk me through a recent time you reviewed a contribution you suspected was AI assisted. What was that like?
- How much of your current strain, if any, do you attribute to GenAI related changes versus everything else?
- Is there a specific moment that stands out as particularly stressful in this space?

Probe: this can be a hard thing to say out loud. If someone hesitates, reiterate that this is confidential and anonymized in any writeup.

Perfect world and desired change (about 8 minutes)

- In a perfect world, what would contributions to your project look like, given these tools exist?
- What is the one thing that, if changed, would most reduce your burden right now?
- What would you need, from tools, from the community, or from other maintainers, to feel less strained?
- If you could send one message to every contributor using GenAI to submit to your project, what would it be?

This is a natural stopping point if time is short.

Extension block (about 30 more minutes, time permitting)

AI-generated contribution quality (about 7 to 8 minutes)

- How do you currently tell, or guess, when a contribution was AI generated or AI assisted?
- How has the volume of AI-influenced contributions changed compared to a year or two ago?
- Does your triage process differ for these versus other contributions? Walk me through it.

- Are there specific incidents with low quality or hallucinated AI-generated submissions that stick with you?

Tool-specific deep dive (about 5 to 6 minutes)

- Which specific tools do you personally use? For example Copilot, ChatGPT, Claude Code, Cursor, or review bots like CodeRabbit.
- Walk me through a specific, recent example of using one of these in your maintainer work.
- Is there anything you tried and abandoned? What did not work?

Before and after comparison (about 5 to 6 minutes)

- Think back to your workflow two or three years ago versus today. What is concretely different?
- Has the kind of work you do changed? For example, less writing code, more reviewing and verifying.

Community and behavior change (about 5 to 6 minutes)

- What would help maintainers collectively, not just you personally? Think policies, norms, tooling.
- Are there other projects or communities you think are handling this well?
- What do you wish the broader open source or pyOpenSci community understood about this that is not well understood yet?

Closing (about 2 to 3 minutes)

Read or paraphrase:

"Is there anything I have not asked about that you think is important here?"

Thank you so much for your candor and your time, this is genuinely useful. If you are interested, there is a survey and a focus group you can join, and I can send you the link. I will follow up once this is written up if you would like to see where it lands."